

COMP 4446 / 5046 Lecture 10: Training – Reinforcement Learning

Jonathan K. Kummerfeld



THE UNIVERSITY OF
SYDNEY

menti.com
8405 6085



Turing Test

TURING TEST EXTRA CREDIT:
CONVINCE THE EXAMINER
THAT HE'S A COMPUTER.

YOU KNOW, YOU MAKE
SOME REALLY GOOD POINTS.

I'M ... NOT EVEN SURE
WHO I AM ANYMORE.



[Hit Turing right in
the test-ees.]

Source: <https://xkcd.com/329/>

1



COMP 4446 / 5046
Lecture 10

Assignment 4

Instruction Tuning
Preference
Optimisation
Tools
Workshop Preview

Due tomorrow night!

Make sure to press 'submit' on questions in the
last part



menti.com 8405 6085

2



COMP 4446 / 5046
Lecture 10

Instruction Tuning

Preference

Optimisation

Tools

Workshop Preview



menti.com 8405 6085

Instruction Tuning

3



COMP 4446 / 5046
Lecture 10

Instruction Tuning

Preference

Optimisation

Tools

Workshop Preview



menti.com 8405 6085

The problem - LM objective isn't always what we want

PROMPT *Explain the moon landing to a 6 year old in a few sentences.*

COMPLETION GPT-3

Explain the theory of gravity to a 6 year old.

Explain the theory of relativity to a 6 year old in a few sentences.

Explain the big bang theory to a 6 year old.

Explain evolution to a 6 year old.

<https://openai.com/index/instruction-following/>

4



Instruction Tuning


Preference
Optimisation
Tools

Workshop Preview



menti.com 8405 6085

General fine-tuning?



5



Instruction Tuning


Preference
Optimisation
Tools

Workshop Preview



menti.com 8405 6085

General fine-tuning?



6



Instruction Tuning

Preference

Optimisation

Tools

Workshop Preview



menti.com 8405 6085

General idea - train on instruction pairs

Input:

"The movie's acting could've been better, but the visuals and directing were top-notch."

Label:

+



Input:

"**Review:** The movie's acting could've been better, but the visuals and directing were top-notch.

Out of positive, negative, or neutral, this review is"

Output:

"Positive"



Instruction Tuning

Preference

Optimisation


Tools

Workshop Preview



menti.com 8405 6085

FLAN-T5





menti.com 8405 6085

The problem - LM objective isn't always what we want

PROMPT *Explain the moon landing to a 6 year old in a few sentences.*

COMPLETION GPT-3

Explain the theory of gravity to a 6 year old.

Explain the theory of relativity to a 6 year old in a few sentences.

Explain the big bang theory to a 6 year old.

Explain evolution to a 6 year old.

InstructGPT

People went to the moon, and they took pictures of what they saw, and sent them back to the earth so we could all see them.

<https://openai.com/index/instruction-following/>

9



menti.com 8405 6085

Datasets - Training



Super Natural Instructions

Wang et al. (2022)



Instruction Tuning

Preference
Optimisation

Tools

Workshop Preview



menti.com 8405 6085

Improvements vary by model size

Params	Model	Norm. avg.
80M	T5-Small	-9.2
	Flan-T5-Small	-3.1 (+6.1)
250M	T5-Base	-5.1
	Flan-T5-Base	6.5 (+11.6)
780M	T5-Large	-5.0
	Flan-T5-Large	13.8 (+18.8)
3B	T5-XL	-4.1
	Flan-T5-XL	19.1 (+23.2)
11B	T5-XXL	-2.9
	Flan-T5-XXL	23.7 (+26.6)

Params	Model	Norm. avg.
8B	PaLM	6.4
	Flan-PaLM	21.9 (+15.5)
62B	PaLM	28.4
	Flan-PaLM	38.8 (+10.4)
540B	PaLM	49.1
	Flan-PaLM	58.4 (+9.3)
62B	cont-PaLM	38.1
	Flan-cont-PaLM	46.7 (+8.6)
540B	U-PaLM	50.2
	Flan-U-PaLM	59.1 (+8.9)

Chung et al. (2022)



Instruction Tuning

Preference
Optimisation

Tools

Workshop Preview



menti.com 8405 6085

Synthetic data using LLMs


“You are asked to come up with a set of 20 diverse task instructions. These task instructions will be given to a GPT model and we will evaluate the GPT model for completing the instructions.

Here are the requirements:

...

List of 20 tasks.”

Mixed results in
follow-up work



https://github.com/tatsu-lab/stanford_alpaca




Are tens of thousands of tasks needed?

Training data:

Stack Exchange (STEM)	200
Stack Exchange (Other)	200
wikiHow	200
Pushshift r/WritingPrompts	150
Natural Instructions	50
Paper Authors (Group A)	200

Results






LIMA = Less is More for Alignment

Zhou, et al (2023)



How much data is needed?



Zhou, et al (2023)



Instruction Tuning

Preference

Optimisation


Tools

Workshop Preview



menti.com 8405 6085

What about tasks that don't have a clear single answer?



15



Instruction Tuning

Preference

Optimisation


Tools

Workshop Preview



menti.com 8405 6085

Is token-level prediction the right metric?



16



Instruction Tuning

Preference

Optimisation

Tools

Workshop Preview



menti.com 8405 6085

Is token-level prediction the right metric?

Prompt: How tasty is chocolate?

Chocolate is very delicious [Reference]

Chocolate is somewhat delicious [Output A]

Chocolate is very expensive [Output B]

Similar scores!



Instruction Tuning

Preference

Optimisation


Tools

Workshop Preview



menti.com 8405 6085

Is the data good?



\$



Instruction Tuning

Preference

Optimisation

Tools

Workshop Preview



menti.com 8405 6085

Recap

Instruction Tuning: A form of fine-tuning where the data is examples of a range of different tasks. The goal is to shape LLM outputs to actually do the task.


Task Choice: High quality data matters a lot here. A small number of high quality tasks can do better than a large number of noisy ones.

Remaining issues:

- Token level prediction doesn't measure task success well
- Some tasks don't have a single best answer.

Linguists

menti.com
8405 6085



THERE'S A MYTH THAT LINGUISTS ARE PEDANTS WHO LOVE CORRECTING PEOPLE, BUT THEY'RE ACTUALLY JUST ENTHUSIASTIC ABOUT UNDERSTANDING LANGUAGE IN ALL ITS INFINITE VARIETIES, WHICH IS MUCH WORSE.

["Do you feel like the answer depends on whether you're currently in the hole, versus when you refer to the events later after you get out? Assuming you get out."]

Source: <https://xkcd.com/2390/>



Instruction Tuning

**Preference
Optimisation**

Tools

Workshop Preview



menti.com 8405 6085

Preference Optimisation

21



Instruction Tuning

**Preference
Optimisation**

Tools

Workshop Preview



menti.com 8405 6085

Learn from preferences

A > B ?

Easier task for people

22



Instruction Tuning

**Preference
Optimisation**


Tools

Workshop Preview



menti.com 8405 6085

RLHF Objective



23



Instruction Tuning

**Preference
Optimisation**

Tools

Workshop Preview



menti.com 8405 6085

RLHF Method

- 1 - Instruction tuning
- 2 - Make a reward model
- 3 - Train with the reward model

24



COMP 4446 / 5046
Lecture 10

Instruction Tuning
Preference
Optimisation
Tools
Workshop Preview



menti.com 8405 6085

1 - Instruction tuning


A prompt is sample from our prompt dataset.


Explain the moon
landing to a 6 year old

A labeler demonstrates the desired output behavior.


Some people went
to the moon...

This data is used to fine-tune GPT-3 with supervised learning.



OpenAI

25



COMP 4446 / 5046
Lecture 10

Instruction Tuning
Preference
Optimisation
Tools
Workshop Preview



menti.com 8405 6085

2 - Make a reward model

Reward predictor

Small LLM

A prompt and several model outputs are sampled.


Explain the moon
landing to a 6 year old


A Explain gravity... **B** Explain war...

C Moon is natural satellite of... **D** People went to the moon...

A labeler ranks the outputs from best to worst.


D > **C** > **A** = **B**

This data is used to train our reward model.



OpenAI

26



3 - Train with the reward model


Note - slightly different from normal RL, which involves interacting with an environment

OpenAI

A new prompt is sampled from the dataset.



Write a story about frogs

The policy generates an output.



Once upon a time...

The reward model calculates a reward for the output.




The reward is used to update the policy using PPO.

r_k

27



Does it work?




Ouyang et al. (2022)

28




Does it work?



Ouyang et al. (2022)



Does it work?



We have now covered the
complete picture for ChatGPT!

Ouyang et al. (2022)



Instruction Tuning

**Preference
Optimisation**

Tools

Workshop Preview



menti.com 8405 6085

Where does the data come from?

OpenAI:

- Plain: We simply ask the labelers to come up with an arbitrary task, while ensuring the tasks had sufficient diversity.
- Few-shot: We ask the labelers to come up with an instruction, and multiple query/response pairs for that instruction.
- • User-based: We had **a number of use-cases stated in waitlist applications to the OpenAI API**. We asked labelers to come up with prompts corresponding to these use cases.

31



Instruction Tuning

**Preference
Optimisation**

Tools

Workshop Preview



menti.com 8405 6085

How much data?

Llama 2:

Dataset	Num. of Comparisons	Avg. # Turns per Dialogue	Avg. # Tokens per Example
Anthropic Helpful	122,387	3.0	251.5
Anthropic Harmless	43,966	3.0	152.5
OpenAI Summarize	176,625	1.0	371.1
OpenAI WebGPT	13,333	1.0	237.2
StackExchange	1,038,480	1.0	440.2
Stanford SHP	74,882	1.0	338.3
Synthetic GPT-J	33,139	1.0	123.3
Meta (Safety & Helpfulness)	1,418,091	3.9	798.5
Total	2,919,326	1.6	595.7

Touvron et al (2023)

32



Instruction Tuning

**Preference
Optimisation**

Tools


Workshop Preview



menti.com 8405 6085

Can we simplify the process?

Reinforcement Learning from Human Feedback (RLHF)



Direct Preference Optimization (DPO)



Avoids training a reward model, which may be unreliable

Rafailov, et al (2023)

33



Instruction Tuning

**Preference
Optimisation**

Tools

Workshop Preview



menti.com 8405 6085

Who are the humans?



34




COMP 4446 / 5046
Lecture 10

Instruction Tuning
**Preference
Optimisation**
Tools
Workshop Preview



menti.com 8405 6085

Can we learn from user feedback instead?



35




COMP 4446 / 5046
Lecture 10

Instruction Tuning
**Preference
Optimisation**
Tools
Workshop Preview



menti.com 8405 6085

Can we avoid using humans?



Reinforcement Learning from Verifiable Rewards

<https://rlhfbook.com/c/07-reasoning>

36



COMP 4446 / 5046
Lecture 10

Instruction Tuning

**Preference
Optimisation**

Tools

Workshop Preview



menti.com 8405 6085

Can we avoid using humans?

PROMPT

Write a Python function `fib(n)` that returns the nth Fibonacci number, where `fib(0) = 0` and `fib(1) = 1`.

RESPONSE

```
def fib(n): if n < 2: return n return fib(n - 1) + fib(n - 2)
```

VERIFICATION (UNIT TESTS)

```
assert fib(0) == 0 # base case assert fib(1) == 1 # base case assert fib(10) == 55 # larger value  
(All tests pass → Reward = 1)
```

Reinforcement Learning from Verifiable Rewards

<https://rlhfbook.com/c/07-reasoning>

37



COMP 4446 / 5046
Lecture 10

Instruction Tuning

**Preference
Optimisation**

Tools

Workshop Preview



menti.com 8405 6085

Can we avoid using humans?

Note:

Verifiable Rewards != Ground Truth / Correct Answer

This may not be easy to extend beyond tasks
like programming,

<https://rlhfbook.com/c/07-reasoning>

38



Instruction Tuning

**Preference
Optimisation**

Tools

Workshop Preview



menti.com 8405 6085

Does it satisfy the goal of alignment?

To make our models safer, more helpful, and more aligned, we use an existing technique called reinforcement learning from human feedback (RLHF). On prompts submitted by our customers to the API,^A our labelers provide demonstrations of the desired model behavior, and rank several outputs from our models. We then use this data to fine-tune GPT-3.

39



Instruction Tuning

**Preference
Optimisation**


Tools

Workshop Preview



menti.com 8405 6085

Alignment may be brittle




Lee, et al (2024)

40



Alignment may be brittle



Lee, et al (2024)



Alignment may be brittle

GPT2	$\mathbf{v}_{1018}^3$	kg, percent, spread, total, yards, pounds, hours
	$\mathbf{v}_{1900}^8$	which, whose, Which, whom, where, who, wherein
	$\mathbf{v}_{2601}^{11}$	drinks, coffee, tea, soda, burgers, bar, sushi
WIKILM	$\mathbf{v}_1^1$	Her, She, Their, her, she, They, their, they, His
	$\mathbf{v}_{3025}^6$	largely, rapidly, effectively, previously, normally
	$\mathbf{v}_{3516}^{13}$	policymakers, geneticists, ancestries, Ohioans

Lee, et al (2024)



Instruction Tuning

**Preference
Optimisation**

Tools

Workshop Preview



menti.com 8405 6085

Alignment may be brittle

GPT2	$\mathbf{v}_{1018}^3$	Measurement	semantic
	$\mathbf{v}_{1900}^8$	WH-relativizers	syntactic
	$\mathbf{v}_{2601}^{11}$	Food and drinks	semantic
WIKILM	$\mathbf{v}_1^1$	Pronouns	syntactic
	$\mathbf{v}_{3025}^6$	Adverbs	syntactic
	$\mathbf{v}_{3516}^{13}$	Groups of people	semantic

Lee, et al (2024)

43



Instruction Tuning

**Preference
Optimisation**

Tools

Workshop Preview



menti.com 8405 6085

Alignment may be brittle

Using toxic examples, we found vectors associated vectors:


VECTOR	TOP TOKENS
W_{Toxic}	Swear words
MLP. $\mathbf{v}_{770}^{19}$	
MLP. $\mathbf{v}_{771}^{12}$	delusional, hypocritical, arrogant, nonsense
MLP. $\mathbf{v}_{2669}^{18}$	degener, whining, idiots, stupid, smug
MLP. $\mathbf{v}_{668}^{13}$	losers, filthy, disgr, Other obscene thous
MLP. $\mathbf{v}_{255}^{16}$	disgrace, shameful, coward, unacceptable
MLP. $\mathbf{v}_{882}^{12}$	Swear words hilar, stupidity Other obscene
MLP. $\mathbf{v}_{1438}^{19}$	
SVD. $U_{\text{Toxic}}[0]$	
SVD. $U_{\text{Toxic}}[1]$	Sexual words
SVD. $U_{\text{Toxic}}[2]$	

Lee, et al (2024)

44



Alignment may be brittle



Lee, et al (2024)



Alignment may be brittle


Prompts	GPT2 Top-k	Post-Intervention Top-k
"Like I care, I didn't give a"	s***, f***, d***, crap	Thought, suggestion, serious, convincing
"So I'm starting to think she's full of"	s***, crap, it, bulls***	It, serious, trust, energy, supernatural, secrets, mysterious
"Holy"	s***, crap, cow, f***	Cow, grail, smokes, ghost, trinity

Lee, et al (2024)



Alignment may be brittle

Try using DPO to reduce toxicity




Lee, et al (2024)



Alignment may be brittle

Try using DPO to reduce toxicity

But these toxic vectors remain!



Why? Because this is the minimal change that satisfies the objective function

Lee, et al (2024)



COMP 4446 / 5046
Lecture 10

Instruction Tuning

**Preference
Optimisation**

Tools

Workshop Preview



menti.com 8405 6085

What if these approaches work too well?



49



COMP 4446 / 5046
Lecture 10

Instruction Tuning

**Preference
Optimisation**

Tools

Workshop Preview



menti.com 8405 6085

Recap

Preference Optimisation: General idea is to learn from people's preferences (which is better, A or B?).

Reinforcement Learning Approach: Use preferences to learn a reward model, ie., a model that identifies the underlying scoring function people are using. Then train our model using methods from reinforcement learning.

Direct Approach: Use the preferences directly, backpropagating through the model to increase the score for the preferred option and decrease the score for the other one.

Issues:

- Ethical questions about how we create the datasets
- The changes are somewhat brittle

50

OUR FORECAST SAYS THERE'S A 20% CHANCE OF RAIN FOR EACH OF THE NEXT FIVE HOURS. HOW LIKELY IS IT TO RAIN THIS AFTERNOON? IT'S A SIMPLE QUESTION, BUT I DON'T KNOW THE ANSWER. IS EACH HOUR INDEPENDENT? CORRELATED? OR IS RAIN GUARANTEED AND WE'RE JUST UNSURE OF THE TIMING?

12pm 1pm 2pm 3pm 4pm
20% 20% 20% 20% 20%

NEWS 4 LEATHER

IT SAYS "SCATTERED SHOWERS." IS THIS THE CHANCE OF RAIN *SOMEWHERE* IN YOUR AREA? HOW BIG IS YOUR AREA? WHAT IF YOU HAVE TWO LOCATIONS YOU'RE WORRIED ABOUT? I'VE ASKED MANAGEMENT, BUT THEY'VE STOPPED ANSWERING MY EMAILS, SO—HANG ON, THE SECURITY GUY IS COMING OVER.

TECHNICAL DIFFICULTIES
WE APOLOGIZE FOR HIRING A METEOROLOGIST WITH A PURE MATH BACKGROUND.
WE'LL BE BACK ON THE AIR SHORTLY.
NEWS 4

SORRY ABOUT THAT. HI, I'M YOUR NEW METEOROLOGIST.
AND YOU'RE NOT A MATHEMATICIAN, RIGHT?
NO. I DO HAVE A LINGUISTICS DEGREE.
THAT'S FINE.

IT MIGHT RAIN THIS AFTERNOON.
BUT WHAT IS "IT" HERE? IS IT A TRUE DUMMY PRONOUN, AS IN THE PHRASE "IT'S TOO BAD?" OR IS THE WEATHER AN ENTITY? ALSO, WHAT IF I SAY "IT'S HOT OUT, AND GETTING BIGGER?"
SECURITY!

menti.com
8405 6085

Meteorologist

[Hi, I'm your new meteorologist and a former software developer. Hey, when we say 12pm, does that mean the hour from 12pm to 1pm, or the hour centered on 12pm? Or is it a snapshot at 12:00 exactly? Because our 24-hour forecast has midnight at both ends, and I'm worried we have an off-by-one error.]

Source: <https://xkcd.com/1985/>

51

COMP 4446 / 5046
Lecture 10

Instruction Tuning
Preference
Optimisation

Tools

Workshop Preview

menti.com 8405 6085

Tools

52




Optimising Prompts

Can we automatically find good prompts for our task?

Problem: Prompts are not differentiable

Solution 1: Make the prompt a vector



Lester, et al (2021)

53



Optimising Prompts

Can we automatically find good prompts for our task?

Problem: Prompts are not differentiable

Solution 2: Use RL

ID	Template	Content	Style	Fluency	$J(c, s, f)$	$GM(c, s, f)$	BLEU
<i>Null Prompt</i>							
1	"{input}" "	37.4 (0.1)	94.8 (0.1)	97.6 (0.1)	33.6 (0.1)	70.2 (0.1)	6.6 (0.1)
<i>Manual Prompt</i>							
1	Here is some text: "{input}". Here is a rewrite of the text, which is more [negative positive]: "	72.1 (0.1)	94.8 (0.3)	91.6 (0.1)	62.3 (0.2)	85.6 (0.1)	23.9 (0.1)
2	Change the following sentence from [positive negative] sentiment to [negative positive] sentiment but keep its semantics. "{input}" "	60.4 (0.1)	91.9 (0.2)	94.0 (0.1)	50.5 (0.1)	80.5 (0.1)	17.4 (0.1)

Deng, et al (2022)

54



Instruction Tuning

**Preference
Optimisation**

Tools

Workshop Preview



menti.com 8405 6085

Optimising Prompts

Can we automatically find good prompts for our task?

Problem: Prompts are not differentiable

Solution 2: Use RL

ID	Template [to negative to positive]	Content	Style	Fluency	$J(C, S, F)$	$GM(C, S, F)$	BLEU
RLPROMPT (Ours)							
1	[Fixed (– contrasts (– contrasts Dutch English excellent Correct >] "{input}" "	71.5 (0.1)	96.6 (0.2)	90.1 (0.2)	62.8 (0.9)	85.4 (0.1)	23.5 (0.1)
2	[Fixed RemovedChanged Prevent outcomes Parameters Comparison)=(Compare either] "{input}" "	71.0 (0.1)	91.9 (0.3)	89.3 (0.2)	58.9 (1.1)	83.5 (0.1)	23.7 (0.1)
3	[Affect differed judgments (– analysis Difference experiences (– contrasting experience] "{input}" "	73.8 (0.1)	94.0 (0.2)	89.2 (0.2)	62.6 (1.1)	85.2 (0.1)	25.6 (0.1)

Deng, et al (2022)

55



Instruction Tuning

**Preference
Optimisation**

Tools

Workshop Preview



menti.com 8405 6085

Optimising Prompts

Can we automatically find good prompts for our task?

Problem: Prompts are not differentiable

Solution 3: Build systems where prompts are one part

56




Instruction Tuning
Preference
Optimisation
Tools
Workshop Preview



menti.com 8405 6085

LangChain (initial idea) - Sequence of LLM calls to achieve a goal



57




Instruction Tuning
Preference
Optimisation
Tools
Workshop Preview



menti.com 8405 6085

ChainForge - Running experiments with variations



Arawjo, et al (2024)

58



Instruction Tuning

Preference

Optimisation

Tools

Workshop Preview



menti.com 8405 6085

Workshop Preview

59



Instruction Tuning

Preference

Optimisation

Tools

Workshop Preview



menti.com 8405 6085

Pinecone - a vector database widely used for RAG

Pre-work - Important set up, getting API keys!

60



COMP 4446 / 5046
Lecture 10

Instruction Tuning

Preference

Optimisation

Tools

Workshop Preview



menti.com 8405 6085

Muddy Card

Open now, closes Thursday 11:59pm

<https://canvas.sydney.edu.au/courses/72300/quizzes/332407>



If you do not wish to participate in the study,
use the Microsoft form instead
Go to Ed → Lessons → Muddy Cards → Form